


APPLYING CNN IN (PNEUMONIA) CHEST IMAGE CLASSIFICATION

TRAN LE



GOAL OF THE PROJECT

Learn CNN

- Understand how CNN works (basically).
- How to run basic CNN with Python

Apply With Chest Image Problem

- Run a simple CNN
- Tune hyperparameters and evaluate the optimal model

OUTLINE

1. Review about
neural network

2. CNN

3. Apply CNN with
the chest image
problem.

- A simple model
- A model with
callbacks
- Tuning
hyperparameter
with and without
callbacks

4. A summarization
about the result




1. NEURAL NETWORK OVERVIEW

1. NEURAL NETWORK

- a) Architecture
- b) Connection between nodes: weights, bias, activation function,
- c) How to find weights and bias: Backpropagation

a) Architecture



Input nodes

Hidden layer

Output nodes


There could be one or more than one hidden layers

1. NEURAL NETWORK

- a) Architecture
- b) Connection between nodes: weights, bias, activation function,
- c) How to find weights and bias: Backpropagation

Input nodes

b) Connection between nodes: weights, bias, activation function.



W: weight

\vec{b} : bias

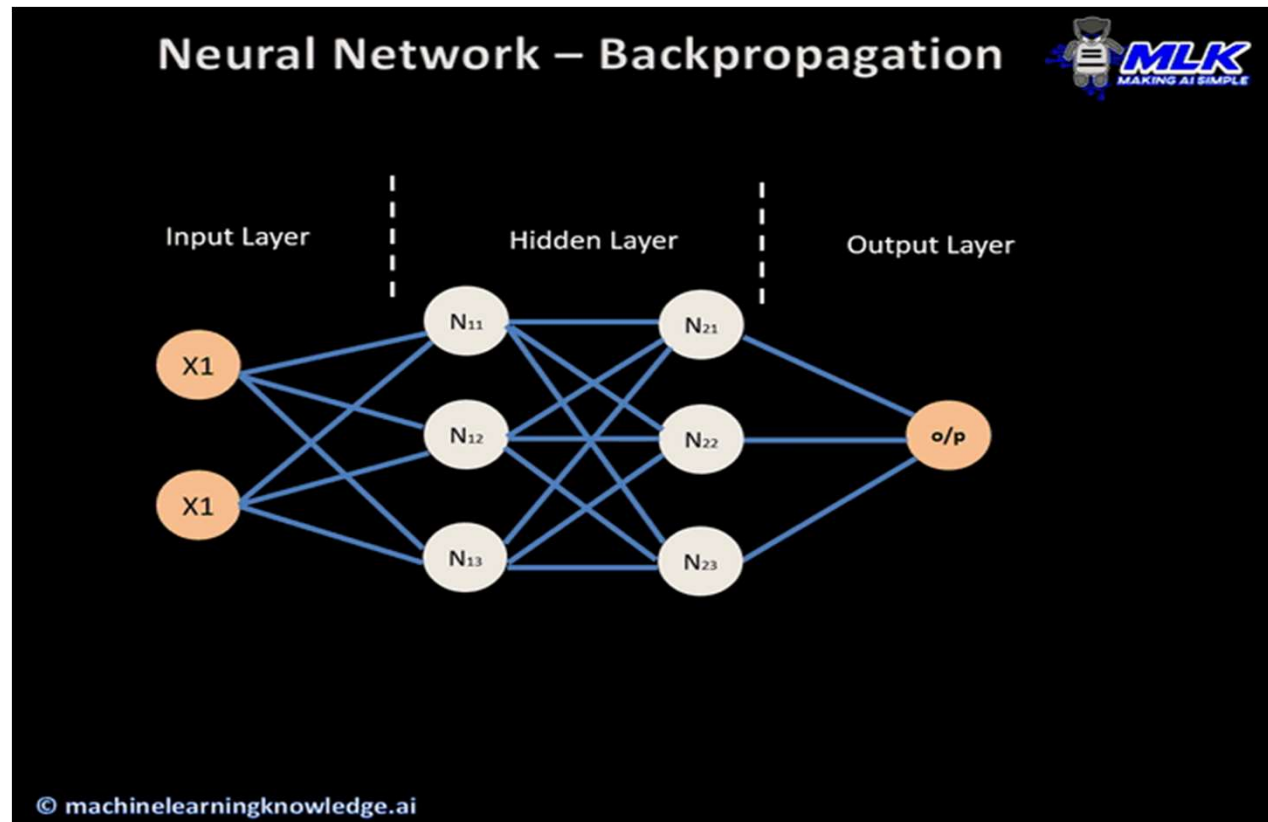
σ : activation function (Linear, RELU, Sigmoid, Softmax, ...). Depending on type of problem to decide the adequate activation function.

1. NEURAL NETWORK

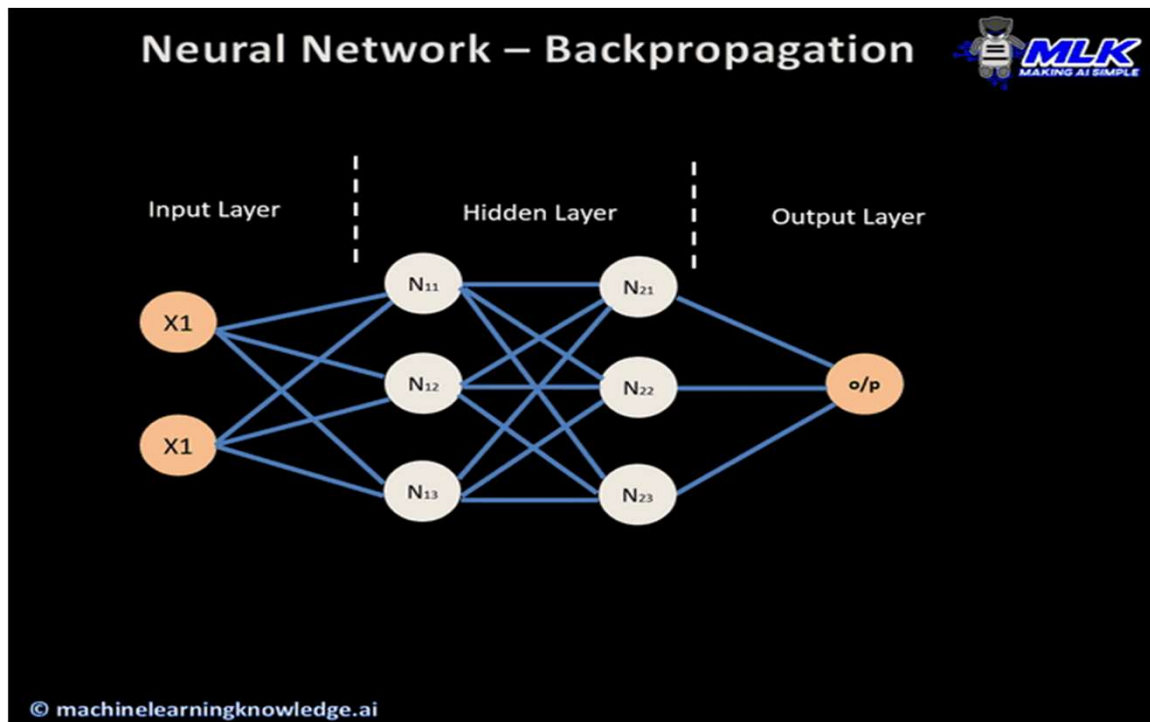
- a) Architecture
- b) Connection between nodes: weights, bias, activation function,
- c) How to find weights and bias: Backpropagation

Input nodes

c) How to find weights and bias: Backpropagation



HOW TO FIND THE WEIGHTS AND BIAS




Epoch: The number of **epochs** is a hyperparameter that defines the number **times** that the learning algorithm will work through the entire training dataset.

Batch: The batch size is a hyperparameter that defines the **number of samples** to work through before updating the internal model parameters.

- **Batch Gradient Descent.** Batch Size = Size of Training Set
- **Stochastic Gradient Descent.** Batch Size = 1
- **Mini-Batch Gradient Descent.** $1 < \text{Batch Size} < \text{Size of Training Set}$

Source: <https://www.google.com/imgres?imgurl=https%3A%2F%2Fmachinelearningknowledge.ai%2Fwp-content%2Fuploads%2F2019%2F10%2FBackpropagation.gif&imgrefurl=https%3A%2F%2Fmachinelearningknowledge.ai%2Fanimated-explanation-of-feed-forward-neural-network-architecture%2F&tbnid=PrPkgqLRXSgkyM&vet=12ahUKewjF082R1JPxAhXqgE4HHW7LChgQMygEegUIARDnAQ..i&docid=sRMKI6I9zyTc3M&w=640&h=480&q=backpropagation%20gif%20&ved=2ahUKewjF082R1JPxAhXqgE4HHW7LChgQMygEegUIARDnAQ>

BATCH

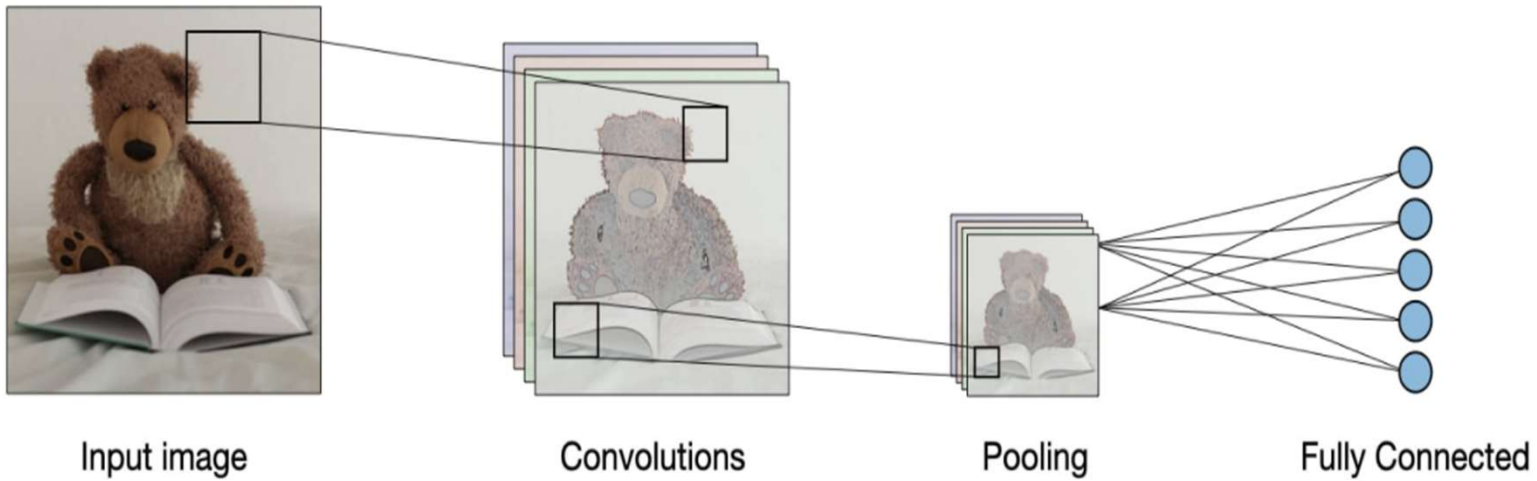


Source: https://www.youtube.com/watch?v=cRd3q4BeRmg&ab_channel=AIQCAR




1. CONVOLUTION NEURAL NETWORK OVERVIEW

CNN ARCHITECTURE



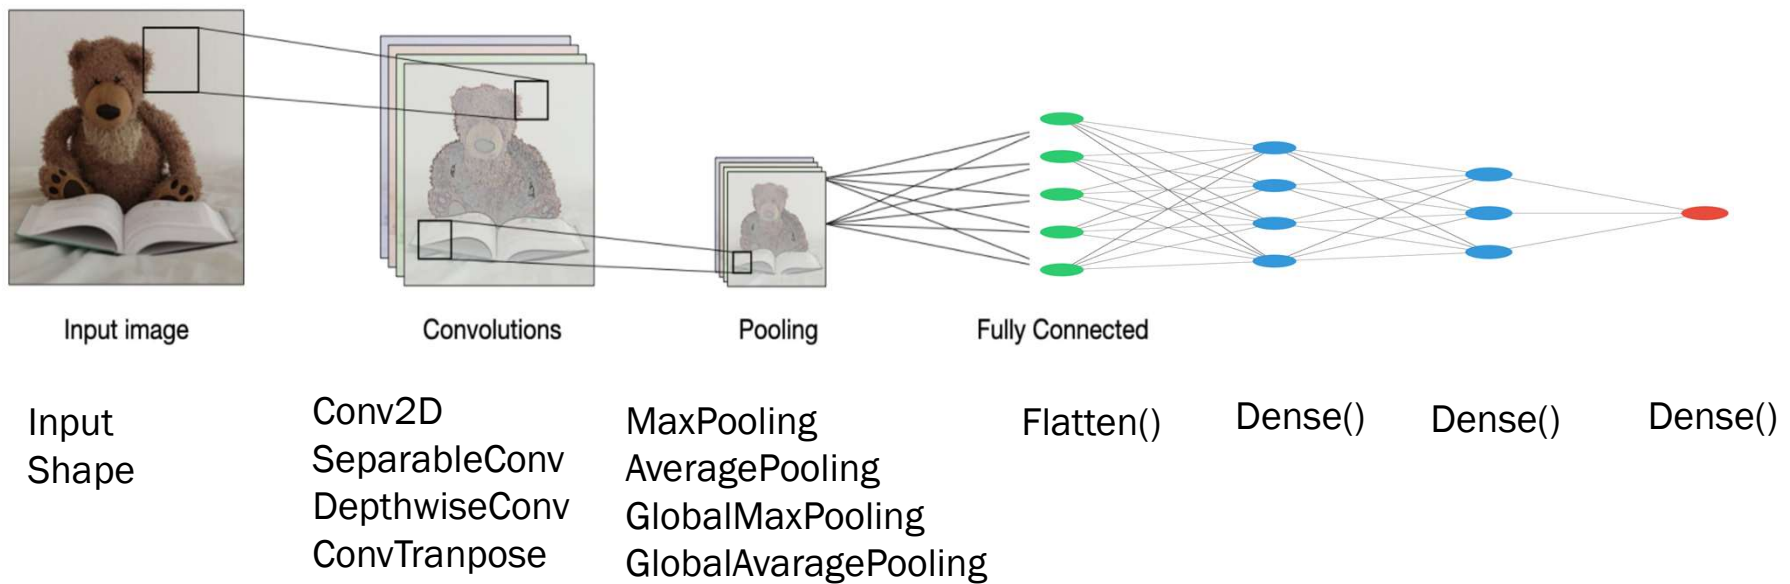
■Source: <https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>

CNN ARCHITECTURE



■ Source: <https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>

CNN ARCHITECTURE





3. APPLYING THE CNN WITH A CHEST IMAGE PROBLEM.

3. A CHEST IMAGE PROBLEM

a. Problem overview

- b. Model 1: A simple model
- c. Model 2: A model with call backs
- d. Model 3: A result from a model tuning hyperparameter
- e. Model 4: A model with tuned hyperparameter without callbacks


a. Problem overview

This is a supervised binary classification problem.

We want to distinguish between normal chest-mages with pneumonia infected chest images

We want to fit the model in training set, and then evaluate on the test set.

Source:
<https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia>



```
Type of the image : <class 'imageio.core.util.Array'>  
Shape of the image : (1858, 2090)  
Dimension of Image 2  
dtype: uint8  
Maximum RGB value in this image 255  
Minimum RGB value in this image 0
```

3. A CHEST IMAGE PROBLEM


- a. Problem overview
- b. Model 1: A simple model
- c. Model 2: A model with call backs
- d. Model 3: A result from a model tuning hyperparameter
- e. Model 4: A model with tuned hyperparameter without callbacks

a. Problem overview (cont.)

The original data sets have 16 observations for validation.

In this project, I change split the train set to get new train and validation sets.

Source:
<https://www.kaggle.com/paultimothymooney/chest-xray-pneumonia>



Frequency of the normal+infected:


	normal	pneumonia
train	1073	3100
validation	268	775
test	234	390

3. A CHEST IMAGE PROBLEM

- a. Problem overview
- b. **Model 1: A simple model**
- c. Model 2: A model with call backs
- d. Model 3: A result from a model tuning hyperparameter
- e. Model 4: A model with tuned hyperparameter without callbacks

b. A simple model

```
model1 = keras.Sequential([
    AveragePooling2D(6,3, input_shape=(img_dims,img_dims, 3)),
    Conv2D(64, 3, activation='relu'),
    Conv2D(32, 3, activation='relu'),
    MaxPool2D(pool_size=(2,2)), # so we use max pool to reduce
    Dropout(0.5), # we can drop out some connection
    Flatten(),
    Dense(128, activation='relu'), #add another dense layer b
    ↪ output layer
    Dense(1, activation='sigmoid')
])
```



```
19/19 [=====] - 9s 476ms/step - loss: 0.5184 -
accuracy: 0.7796
[0.5184427499771118, 0.7796052694320679]
Accuracy = 0.7796052694320679
```

SOME PROBLEMS TO BE CONSIDERED

- Vanishing/Exploding Gradient Problems
 - Use Call back
 - Overfitting
 - Hyperparameter tuning
- Vanishing/Exploding Gradient Problems: the gradients/signals die out or explode and saturate.
 - Glorot and He Initialization:
 - Nonsaturating Activation Functions: *leaky ReLU*.
 - Batch Normalization

SOME PROBLEMS TO BE CONSIDERED

- Vanishing/Exploding Gradient Problems
- Use Call back
- Overfitting
- Hyperparameter tuning

- What is callback?
 - Command callbacks :
 - ModelCheckpoint: helps save checkpoint (by default) at the end of each epoch . Helps return the best model on the validation set.
 - EarlyStopping: will interrupt training when it measures no progress on the validation set for a number of epochs and it will optionally roll back to the best model.
 - Why use callbacks?
 - Avoid overfitting the training set.
 - Avoid wasting time and resources):

SOME PROBLEMS TO BE CONSIDERED

- Vanishing/Exploding Gradient Problems
- Use Call back
- Overfitting
- Hyperparameter tuning

Overfitting

- **1. Simplifying The Model:** decrease the complexity of the model by remove layers or reduce the number of neurons.
- **2. Early Stopping:**
- **3. Use Data Augmentation:**
- **4. Use Regularization (L1, L2, Max-Norm regularization)**
- **5. Use Dropouts (with fixed dropout rate) and Monte-Carlo**

SOME PROBLEMS TO BE CONSIDERED

- Vanishing/Exploding Gradient Problems
- Use Call back
- Overfitting
- Hyperparameter tuning

Overfitting

- **1. Simplifying The Model:** decrease the complexity of the model by remove layers or reduce the number of neurons.
- **2. Early Stopping:**
- **3. Use Data Augmentation:**
- **4. Use Regularization (L1, L2, Max-Norm regularization)**
- **5. Use Dropouts (with fixed dropout rate) and Monte-Carlo**

3. A CHEST IMAGE PROBLEM

- a. Problem overview
- b. Model 1: A simple model
- c. **Model 2: A model with call backs**
- d. Model 3: A result from a model tuning hyperparameter
- e. Model 4: A model with tuned hyperparameter without callbacks

c. A model with call back

```
model2 = keras.Sequential([
    Conv2D(filters=16, kernel_size=(3, 3), activation='relu', padding='same',
    ↪input_shape=[img_dims, img_dims, 3]),
    BatchNormalization(), # use Batch Normalization to address the vanishing/
    ↪exploding gradients problems
    MaxPool2D(pool_size=(2, 2)),

    Conv2D(64, 3, activation='relu'),
    Conv2D(32, 3, activation='relu'),

    #Note: Instead of using a convolutional layer with a 5 x 5 kernel,
    # it is generally preferable to stack two layers with 3 x 3 kernels:
    # it will use less parameters and require less computations, and will
    ↪usually perform better (page447 Hand-on)
    SeparableConv2D(filters=64, kernel_size=(3, 3), activation='relu',
    ↪padding='same'),
    SeparableConv2D(filters=64, kernel_size=(3, 3), activation='relu',
    ↪padding='same'),
    BatchNormalization(),
    MaxPool2D(pool_size=(2, 2)),

    #Note: Note that the number of filters grows as we climb up the CNN towards
    ↪the output layer (64, 128, 256) (page448 Hanh-on)
    SeparableConv2D(filters=128, kernel_size=(3, 3), activation='relu',
    ↪padding='same'),
    SeparableConv2D(filters=128, kernel_size=(3, 3), activation='relu',
    ↪padding='same'),
    BatchNormalization(),
    MaxPool2D(pool_size=(2, 2)),
    Dropout(0.5),
    Dense(128, activation='relu'), #add another dense layer before doing the
    ↪output layer
    Dropout(rate=0.5), #Note: dropout to prevent overfitting
    Dense(units=64, activation='relu'),
    Dropout(rate=0.5),
    Dense(1, activation='sigmoid')
])
```

```
SeparableConv2D(filters=256, kernel_size=(3, 3), activation='relu',
    ↪padding='same'),
    SeparableConv2D(filters=256, kernel_size=(3, 3), activation='relu',
    ↪padding='same'),
    BatchNormalization(),
    MaxPool2D(pool_size=(2, 2)),
    Dropout(0.2),

    #Note: we must flatten its inputs, since a dense network expects a 1D array
    ↪of features for each instance
    Flatten(),
    Dense(128, activation='relu'), #add another dense layer before doing the
    ↪output layer
    Dropout(rate=0.5), #Note: dropout to prevent overfitting
    Dense(units=64, activation='relu'),
    Dropout(rate=0.5),
    Dense(1, activation='sigmoid')
])
```

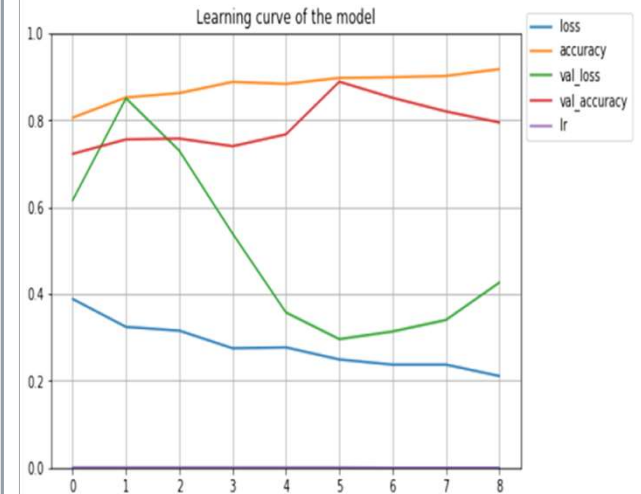
3. A CHEST IMAGE PROBLEM

- a. Problem overview
- b. Model 1: A simple model
- c. **Model 2: A model with call backs**
- d. Model 3: A result from a model tuning hyperparameter
- e. Model 4: A model with tuned hyperparameter without callbacks

c. A model with call back

```
# Callbacks
checkpoint = ModelCheckpoint(filepath='best_weights.hdf5', save_best_only=True,
                             save_weights_only=True)
lr_reduce = ReduceLROnPlateau(monitor='val_loss', factor=0.3, patience=4,
                               verbose=2, mode='max')
early_stop = EarlyStopping(monitor='val_loss', min_delta=0.1, patience=4,
                            mode='min')

epochs=25
model2_hist = model2.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // batch_size,
    epochs=epochs,
    validation_data=valid_generator,
    validation_steps= valid_generator.samples // batch_size,
    callbacks=[checkpoint, lr_reduce, early_stop])
```



19/19 [=====] - 4s 191ms/step - loss: 0.2470 -
accuracy: 0.9013
[0.24703054130077362, 0.9013158082962036]
Accuracy = 0.9013158082962036

SOME PROBLEMS TO BE CONSIDERED

- Vanishing/Exploding Gradient Problems
- Use Call back
- Overfitting
- Hyperparameter tuning

- Hyperparameter tuning
 - number of layers,
 - number of neurons per layer,
 - type of activation function to use in each layer,
 - weight initialization
 - Learning rate, batch size
- Use GridSearchCV or RandomizedSearchCV

3. A CHEST IMAGE PROBLEM

- a. Problem overview
- b. Model 1: A simple model
- c. Model 2: A model with call backs
- d. **Model 3: A result from a model tuning hyperparameter**
- e. Model 4: A model with tuned hyperparameter without callbacks

d. A model with tuned hyperparameter with callbacks

```
def build_model_withcallback(hp):  
    """hp: stands for hyper-parameters"""  
    model = keras.Sequential()  
  
    model.add(Conv2D(filters=16, kernel_size=(3, 3), activation='relu',  
padding='same', input_shape=[img_dims, img_dims, 3]))  
    model.add(BatchNormalization()) # use Batch Normalization to address  
the vanishing/exploding gradients problems  
    model.add(MaxPool2D(pool_size=(2, 2)))  
  
    for i in range(hp.Int("Conv Layers", min_value=0, max_value=2)):  
  
        model.add(keras.layers.Conv2D(hp.Choice(f"layer_{i}_filters",  
[64,128,256]), 3, activation='relu', padding='same'))  
        model.add(keras.layers.MaxPool2D(2,2))  
        model.add(keras.layers.Dropout(0.5))  
  
        for i in range(hp.Int("Conv Layers", min_value=0, max_value=2)):  
            model.add(keras.layers.Conv2D(hp.Choice(f"layer_{i}_filters",  
[64,128,256]), 3, activation='relu', padding='same'))  
            model.add(keras.layers.MaxPool2D(2,2))  
            model.add(keras.layers.Dropout(0.5))  
  
        # Flatten  
        model.add(keras.layers.Flatten())  
        model.add(keras.layers.Dense(hp.Choice("Dense layer", [64, 128, 256]),  
activation='relu'))  
        model.add(keras.layers.Dense(1, activation='sigmoid'))  
        # Compile  
        model.compile(optimizer='adam',  
                        loss='binary_crossentropy',  
                        metrics=['accuracy'])  
    return model
```

```
tuner_withcallback = RandomSearch(  
    build_model_withcallback,  
    objective='val_accuracy',  
    max_trials=20,  
)  
  
tuner_withcallback.search(train_generator,  
    steps_per_epoch=train_generator.samples//batch_size,  
    validation_data=valid_generator,  
    epochs=15,  
    batch_size=batch_size,  
    callbacks=[tf.keras.callbacks.EarlyStopping(patience=2)])
```

3. A CHEST IMAGE PROBLEM

- a. Problem overview
- b. Model 1: A simple model
- c. Model 2: A model with call backs
- d. **Model 3: A result from a model tuning hyperparameter**
- e. Model 4: A model with tuned hyperparameter without callbacks

d. A model with tuned hyperparameter with callbacks

```
Best val_accuracy So Far: 0.923298180103302
```

```
Total elapsed time: 05h 10m 37s
```

```
INFO:tensorflow:Oracle triggered exit
```

```
19/19 [=====] - 8s 406ms/step - loss: 0.5907 -
```

```
accuracy: 0.8026
```

```
[0.5906840562820435, 0.8026315569877625]
```

3. A CHEST IMAGE PROBLEM

- a. Problem overview
- b. Model 1: A simple model
- c. Model 2: A model with call backs
- d. Model 3: A result from a model tuning hyperparameter
- e. Model 4: A model with tuned hyperparameter without callbacks

a. A model with tuned hyperparameter without callbacks

Best val_accuracy So Far: 0.9587727785110474

Total elapsed time: 12h 46m 41s

INFO:tensorflow:Oracle triggered exit

39/39 [=====] - 16s 392ms/step - loss: 0.4852 - accuracy: 0.8674
[0.3162696361541748, 0.9086538553237915]

SUMMARY THE RESULT:

1. Simple model

- Valid accuracy: 0.8438
- Test accuracy: 0.7796

2. A model with call backs

- Valid accuracy: 0.8887
- Test accuracy: 0.9013

3. A model with tuned hyper-parameter with callbacks

- Best Valid accuracy: 0.9232
- Test accuracy: 0.8026

4. A model with tuned hyper-parameter without callbacks

- Best Valid accuracy: 0.9587
- Test accuracy: 0.9086



SOME PROBLEMS THAT I NEED TO LEARN

- Not use data augmentation to increase the size of training set
- Convolution layer design
- Learning rate
- An approach getting 100% accuracy using Fast AI, a course of Jeremy Howard.

https://www.youtube.com/watch?v=gGxe2mN3kAg&list=RDCMUCX7Y2qWriXpqocG97SFW20Q&start_radio=1&rv=gGxe2mN3kAg&t=0&ab_channel=JeremyHoward%28youtubeaccount%29JeremyHoward%28youtubeaccount%29

- Use Tensor-board to optimize CNN model

REFERENCE

- Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow, 2nd Edition, by Aurélien Géron, Released September 2019
- Deep Learning with Python, Second Edition, François Chollet, MEAP began March 2020
- <http://cs231n.stanford.edu/>
- <https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>
- <https://www.kaggle.com/jagadish13/keras-nn-x-ray-predict-pneumonia>
- [A solution, get 100%, using Fastai] <https://towardsdatascience.com/fastai-bag-of-tricks-experiments-with-a-kaggle-dataset-part-1-135e46da72f2>